
Diagnosis of Faults in Data Centre Networks using Data Plane Programmability and Relational Queries

MID SEMESTER REPORT

*Submitted in partial fulfillment of the requirements of
BITS F421T Thesis*

By

Sankalp Sanjay Sangle
ID No. 2016A7TS0110P

Under the supervision of:

Dr. Mun Choon CHAN
&
Dr. Virendra S SHEKHAWAT



BIRLA INSTITUTE OF TECHNOLOGY AND SCIENCE PILANI, PILANI CAMPUS

March 2020

Acknowledgements

Sup, cosup, Family, Pravein, Nishant, Miloni. . .

Contents

Acknowledgements	i
Contents	ii
List of Figures	iv
List of Tables	v
Abbreviations	vi
1 Introduction to Thesis Topic	1
2 Background Literature	2
2.1 Why do we need programmable networks?	2
2.2 SDN and the idea of separation of Control Plane and Data Plane	3
2.3 A Brief History of Programmability in Computer Networks	4
2.3.1 Early approaches towards programmability (1990s - late 2000s)	4
2.3.2 OpenFlow (and its limitations)	4
2.4 Towards Data Plane Programmability	5
2.4.1 Constraints in data plane programmable switches	6
2.4.2 The P4 Programming Language	6
3 Experiments	7
3.1 Experimental Setup	7
3.2 The Synchronous Incast Problem	9
3.2.1 Description	9
3.2.2 Configuration	9
3.2.3 Diagnosis	9
3.2.4 Results and Illustrations	9
3.3 Asynchronous Incast problem and Heavy Hitters	9
3.3.1 Description	9
3.3.2 Configuration	9
3.3.3 Diagnosis	9
3.3.4 Results and Illustrations	9
3.4 Link Underprovisioning	9
3.4.1 Description	9
3.4.2 Configuration	9
3.4.3 Diagnosis	9

3.4.4	Results and Illustrations	9
3.5	A unified scheme	9
A	Appendix A: Jain's Fairness Index	10

List of Figures

2.1	Match Action Pipeline	5
3.1	Evaluation Topology	7

List of Tables

Abbreviations

LAH List Abbreviations **Here**

Chapter 1

Introduction to Thesis Topic

Data Centers have emerged as the basic unit for providing web services to customers by major technology companies. They house servers in the thousands and tens of thousands, with servers often serving as backups to ensure reliability. Traffic in data centers is often organized in a tree like structure, with top of rack switches, aggregate switches and core switches. The tree like structure is easy to maintain and scale to meet the growing number of clients.

The reliability of data center network infrastructure is critically important for ensuring seamless user experience. Companies like Facebook and Google have spent considerable investments into developing frameworks for automating fault detection and repair of networks. Even so, the nature and distribution of faults in data center networks isn't completely understood, and many faults like microbursts due to fan-in traffic go undetected. A closer look is needed at the type of faults that occur in data center networks and on how they can be diagnosed.

Modern data centres operate at throughputs of close to a few hundreds of Gbps, implying packets coming in at a few hundreds of nanoseconds apart. Most monitoring solutions offer resolution of the order of milliseconds, and hence are inadequate for recording and detection of events that last for sub millisecond intervals. This motivates the need for having a diagnosis system to distinguish events at nanosecond level by leveraging advances in data plane programming.

This thesis aims to discuss techniques for diagnosing faults in data centers at nanosecond resolution. This resolution is provided by existing data plane time synchronization protocols. Data plane switches are used to record network telemetry data in a relational database, which is then queried using SQL to diagnose possible faults in a network and create relevant visualizations which can help a network administrator determine the root cause of faults in the network and address them.

Chapter 2

Background Literature

Historically, computer networks have always been relatively blackboxed entities that provide no substantial facilities for configuration. The devices that are present in them, from traditional switches and routers, to other middleware like firewalls and network address translators, often execute software that is proprietary, and hence rigid by nature. These devices are configured by network administrators using highly specialised interfaces that vary from vendor to vendor. The software for communicating with these interfaces is not centralized, and network devices need to be individually configured to achieve the desired functionality. This frozen structure has hindered flexibility and increased costs of network management.

It is interesting to note that this structure was by design and not an oversight. The nascent stage Internet needed to be absolutely reliable and fault tolerant if widespread adoption was to take place. One way to minimize failures in the network was to make the network devices rigid and inflexible. However, by the turn of the century, the increase in number of end hosts and the widespread adoption of the Internet has given rise to a need for managing large flows of traffic effectively. The end host explosion has made the question of programmable networks much more pragmatic to discuss.

2.1 Why do we need programmable networks?

The ability to remotely configure behaviour of switches/routers finds use in a number of applications, some of them being:

- Dynamic Routing of heavy hitters

- Testing of new protocols/routing policies
- In-Network Computing
- Layer 4 Load Balancing
- Network Monitoring and Debugging (The focus of this thesis)

A programmable network is easier to monitor and control than traditional networks. Network devices whose forwarding policies can be customized can be made to behave like a router, a switch, a firewall, a NAT, or any other device we can conceptualize within the constraints of data plane switch programmability. This saves money and physical labour that would have been spent in the purchase and set up of highly specific-function network devices. A programmable network is, by nature, facilitative to network innovation and reduces the barrier to introduction of new protocols/policies.

2.2 SDN and the idea of separation of Control Plane and Data Plane

Software Defined Networking(SDN) is an umbrella term used to address efforts made to make networks programmable. The behaviour of traditional switches is defined in hardware, whereas 'software-defined' switches would have the programmability addressed earlier. One important concept of SDN is *the decoupling of Control and Data planes*. Control plane refers to protocols like OSPF, BGP, Multicast which govern traffic handling on a higher scale. The data plane performs packet switching based on the policies that the control plane dictates. In a general sense, the control plane is the intelligence layer, and the data plane is the manifestation of it in a standalone switch.

A traditional switch has the forwarding logic baked into the circuit, and tight integration of data and control planes. The idea behind SDN is to decouple the planes, and have a programmable control plane that can communicate with the data plane via a standardized API (like OpenFlow) so as to offer freedom in specifying how a switch handles packets rather than 'hardcoding' it into the switch.

2.3 A Brief History of Programmability in Computer Networks

2.3.1 Early approaches towards programmability (1990s - late 2000s)

The early days of SDN research revolved around what were known as Active Networks. In this, two approaches were used:

- *Capsule model*, where code to be run at the nodes is carried within packets
- *Programmable router/switch model* where the code is pushed to the nodes by out-of-band means

While capsule models were feasible, there were concerns about attackers injecting malicious code into a router and compromising the network. The programmable model was much more aligned to what SDN is today. Most of the opposition towards adopting these approaches was due to a mixture of concerns over reliability and performance, and also because it was important to ensure proliferation of the Internet by keeping the network core as simple as possible.

2.3.2 OpenFlow (and its limitations)

OpenFlow started off as a research project that was implemented at Stanford University by researchers in a campus wide network. As mentioned before, it defines the API for communication between a decoupled data plane and control plane. OpenFlow quickly became popular due to it being easy to adopt (most commodity switches required a firmware upgrade). The OpenFlow specification provides means to specify entries in match-action tables(called rules) where a match is made against the headers of the packet, and an action is taken(forwarding, dropping, broadcasting). Widespread adoption by companies like Google boosted its popularity and many new switch vendors entered the established markets by undertaking early adoption of OpenFlow. There is a false belief that OpenFlow is the same as SDN; SDN is a general term used for addressing techniques that make networks programmable. It provides no direction or proposal on *how* to go about doing that. OpenFlow is merely one concrete way to give a level of programmability to networks.

OpenFlow however, does not offer any data plane programmability support. There is no means to specify new protocols and fields to match against, or on how to perform reassembly of packets; consider a scenario where one needs to test a new protocol. OpenFlow enabled switches will not

be able to perform match-actions on fields specified in the protocol and will have to wait until that protocol gets added into the specification (which can take of the order of months/years). This use case motivates the desire to make the data plane itself programmable.

2.4 Towards Data Plane Programmability

Data plane programming offers the flexibility of describing how the packet headers can be parsed and matched onto non-standard fields (for example, those of a new protocol) and modified, as well as the order in which the headers are reassembled before being transmitted by the switch. It is natural to think that there will be a tradeoff between throughput/speed and customizability, and this was the notion in the research community for a number of years. However, programmable switches have been made possible by recent advances networking architectures in Figure 2.1. They expose such functionality in the data plane:

- Flexible parsing and deparsing on packet headers
- Ingress/Egress processing using match-action tables
- Stateful maintenance of network states using SRAMs

These switches offer throughputs of close to 1 Tbps (the rate demanded by modern data centres) even with the ability to process and modify the packet. Such speeds would be impossible if packet processing was to happen in CPU due to various overheads, which is what makes data plane programmability so attractive.




FIGURE 2.1: Generic Data Programmable Switch Architecture

2.4.1 Constraints in data plane programmable switches

In order to maintain packet processing at line-rate with no slow down in throughput, certain constraints need to be followed which impose a limit on the packet processing flexibility. Some of these are:

- No looping constructs in P4 (the language used to program switches)
- no floating point computations
- no exact multiply/divides (only approximated by bit-shifts)
- only one read-modify-write per stage to maintain processing at line rate

Even so, data plane programmability has found a number of applications in recent years including Network Monitoring and Debugging (the subject matter of this thesis), time synchronization, detection of elephant or heavy hitter flows, and In-Network Computing.

2.4.2 The P4 Programming Language

P4 (*Programming Protocol Independent Packet Processors*) is a programming language developed specifically for programming data plane switches. The language is maintained by the P4 Language Consortium and is widely used for data plane programmability today. It provides constructs for specifying deterministic parsers for header parsing, action constructs for specifying steps to be taken in case of a match, and also steps to be taken in packet reassembly.

Chapter 3

Experiments

3.1 Experimental Setup

To simulate a fat-tree topology as seen in data center networks, 4 physical servers and 2 data plane programmable switches were used. Each switch was virtualized to create 5 switches using 10G loopback links. So, we have 10 switches in total, with 4 of them as top-of-rack or edge switches, 4 of them as aggregate switches, and 2 as core switches (Figure 3.1)




FIGURE 3.1: Evaluation Topology

The rest of this chapter is structured as follows: For a potential fault that can occur in the network, we have

- A section on the description of the fault
- A section on the configuration of switches and hosts for reproducing the fault
- A section on how the proposed fault can be diagnosed using SQL queries
- Results and relevant illustrations

The last section presents a unified scheme for diagnosing faults in network built on logic that has been developed in preceding sections.

3.2 The Synchronous Incast Problem

3.2.1 Description

3.2.2 Configuration

3.2.3 Diagnosis

3.2.4 Results and Illustrations

3.3 Asynchronous Incast problem and Heavy Hitters

3.3.1 Description

3.3.2 Configuration

3.3.3 Diagnosis

3.3.4 Results and Illustrations

3.4 Link Underprovisioning

3.4.1 Description

3.4.2 Configuration

3.4.3 Diagnosis

3.4.4 Results and Illustrations

3.5 A unified sceme

Appendix A

Appendix A: Jain's Fairness Index

Description on the Jain's Fairness Index